

Intro to Monsoon and Slurm

1/18/2017

Introductions

- Introduce yourself
 - Name
 - Department / Group
 - What project(s) do you plan to use monsoon for?
 - Linux or Unix experience
 - Previous cluster experience?

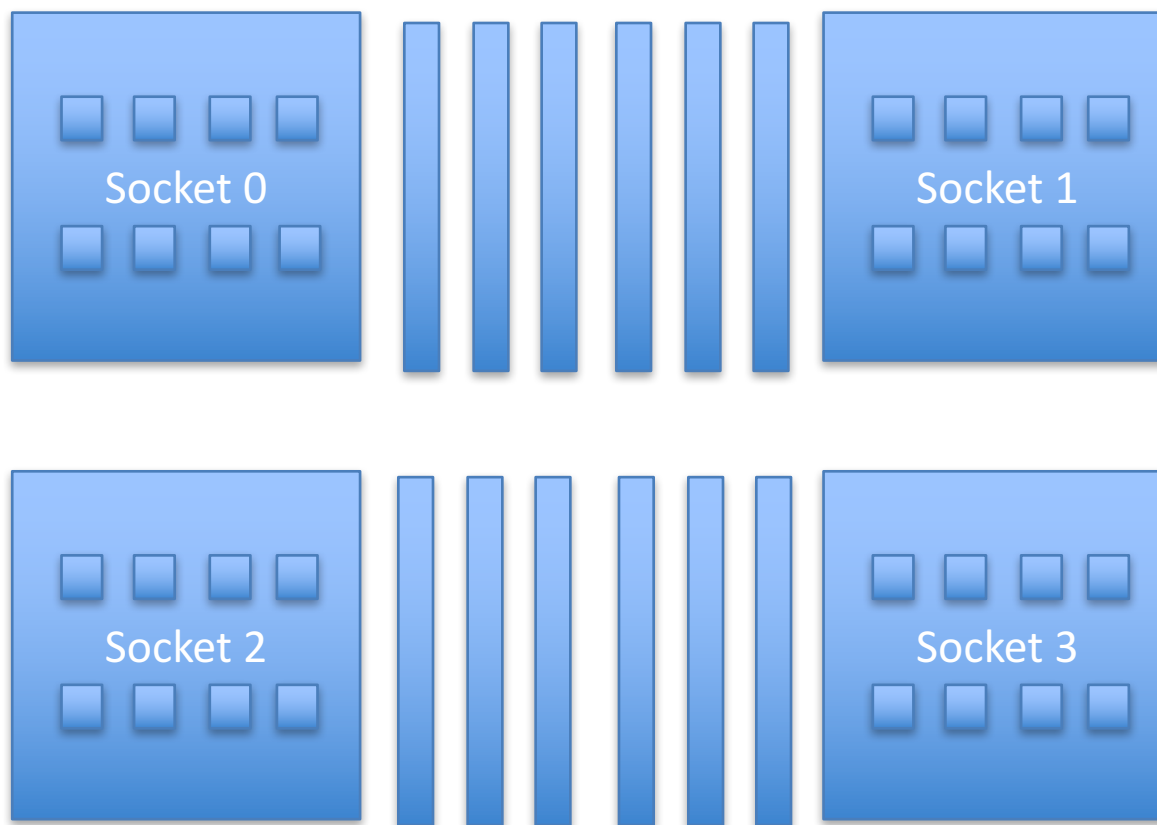
List of Topics

- Cluster education
 - What is a cluster, exactly?
 - Queues, scheduling and resource management
- Cluster Orientation
 - Monsoon cluster specifics
 - How do I use this cluster?
 - Group resource limits
 - Exercises
 - Question and answer

What is a cluster?


- A computer cluster is many individual computers systems (nodes) networked together locally to serve as a single resource
- Ability to solve problems on a large scale not feasible alone

Inside a Node



What is a queue?

- Normally thought of as a line, FIFO
- Queues on a cluster can be as basic as a FIFO, or far more advanced with dynamic priorities taking into consideration many factors



What is scheduling?

- *“A plan or procedure with a goal of completing some objective within some time frame”*
- Scheduling for a cluster at the basic level is much the same. Assigning work to computers to complete objectives within some time availability
- Not exactly that easy though. Many factors come into play scheduling work on a cluster.

Scheduling

- A scheduler needs to know what resources are available on the cluster
- Assignment of work on a cluster is carried out most efficiently with scheduling and resource management working together

Resource Management


- Monitoring resource availability and health
- Allocation of resources
- Execution of resources
- Accounting of resources

Cluster scheduling goals

- Optimize quantity of work
- Optimize usage of resources
- Service all users and projects justly
- Make scheduling decisions transparent

Cluster Resources

- Node
- Memory
- CPU's
- GPU's
- Licenses



Many scheduling methods

- FIFO
 - Simply first in first out
- Backfill
 - Runs smaller jobs with lower resource requirements while larger jobs wait for higher resource requirements to be available
- Fairshare
 - Prioritizes jobs based on users recent resource consumption

Monsoon

- The Monsoon cluster is a resource available to the NAU research enterprise
- 32 systems (nodes) – cn[1-32]
- 884 cores
- 12 GPUs, NVIDIA Tesla K80
- Red Hat Enterprise Linux 6.7
- 12TB memory - 128GB/node min, 1.5TB max
- 170TB high-speed scratch storage
- 500TB long-term storage
- High speed interconnect: FDR Infiniband

Monsoon scheduling

- Slurm (Simple Linux Utility for Resource Management)
- Excellent resource manager and scheduler
- Precise control over resource requests
- Developed at LLNL, continued by SchedMD
- Used everywhere from small clusters to the largest clusters:
 - Sunway TaihuLight (#1), 10.6M cores, 93 PF, 15k KW - China
 - Titan (#3), 561K cores, 17.6 PF, 8k KW, United States

Small Cluster!



Dual core?

Largest Cluster!



10,649,600 cores

Monsoon scheduling

- Combination of scheduling methods
- Currently configured to utilize backfill along with a multifactor priority system to prioritize jobs



Factors attributing to priority

- Fairshare (predominant factor)
 - Priority points determined on users recent resource usage
 - Decay half life over 1 days
- QOS (Quality of Service)
 - Some QOS have higher priority than others, for instance: debug
- Age – how long has the job sat pending
- Job size - the number of nodes/cpus a job is requesting

Storage

- /home – 10GB quota
 - Keep your scripts and executables here
 - Snapshotted twice a day: /home/.snapshot
 - Please do not write job output (logs, results) here!!
- /scratch – 30 day retention
 - Very fast storage, capable of 11GB/sec
 - Checkpoints, logs
 - Keep all temp/intermediate data here
 - Should be your default location to perform input/output

Storage

- /projects
 - Long-term storage project shares
 - Space is assigned to faculty member for group to share
 - Snapshots available
 - No backups today
- /common
 - Cluster support share
 - Contrib: place to put scripts/libs/confs/db's for others use

Data Flow

1. Keep scripts and executables in /home
2. Write temp/intermediate data to /scratch
3. Copy data to /projects/<group_project>, for group storage and reference in other projects
4. Cleanup /scratch

** Remember, /scratch is a scratch filesystem, used for high-speed temporary, and intermediate data

Remote storage access

- scp
 - scp files nauid@monsoon.hpc.nau.edu:/scratch/nauid
 - WinSCP
- samba / cifs
 - \\nau.froot.nau.edu\cirrus (windows)
 - smb://nau.froot.nau.edu/cirrus (mac)

Groups

- NAU has a resource called Enterprise groups
- They are available to you on the cluster if you'd like to manage access to data
- Manage access to your files
- <https://my.nau.edu>
 - “Go to Enterprise Groups”
 - Take a look at our FAQ :: nau.edu/hpc/faq
- If they are not working for you, contact ITS help desk

Software

- ENVI / IDL
- Matlab
- Intel Compilers, and MKL
- R
- Qiime
- Anaconda Python
- OpenFOAM
- SOWFA
- Lots of bioinformatics programs
- Request additional software to be installed!

Modules

- Software environment management handled by the *modules* package management system
- module avail – what modules are available
- module list – modules currently loaded
- module load <module name> - load a package module
- module display <module name> - detailed information including environment variables effected

MPI

- Quick note on MPI
- Message Passing Interface for parallel computing
- Open MPI set as default MPI
- Mvapich2 also available
 - module unload openmpi
 - module load mvapich2
- Example MPI job script:
 - `/common/contrib/examples/job_scripts/mpijob.sh`

Interacting with Slurm

- Decide what you need to accomplish
- What resources are needed?
 - 2 cpus, 12GB memory, for 2 hours?
- What steps are required?
 - Run prog1, then prog2 ... etc
 - Are the steps dependent on one another?
- Can your work, or project be broken up into smaller pieces?
Smaller pieces can make the workload more agile.
- How long should your job run for?
- Is your software multithreaded, uses OpenMP or MPI?

Example Job script

- `#!/bin/bash`
- `#SBATCH --job-name=test`
- `#SBATCH --output=/scratch/nauid/output.txt`
- `#SBATCH --time=20:00` `# shorter time = sooner start`
- `#SBATCH --workdir=/scratch/nauid`

- `# replace this module with software required in your workload`
- `module load python/3.3.4`

- `# example job commands`
- `# each srun command is a job step, so this job will have 2 steps`
- `srun sleep 300`
- `srun python -V`

Interactive / Debug Work

- Run your compiles and testing on the cluster nodes by:
 - `srun -p all gcc hello.c -o a.out`
 - `srun -p all -c12 make -j12`
 - `srun -qos=debug -c12 make -j12`
 - `srun Rscript analysis.r`
 - `srun python analysis.py`
 - `srun cp -av /scratch/NAUID/lots_o_files /scratch-lt/NAUID/destination`

Long Interactive work

- salloc
 - Obtain a SLURM job allocation that you can work with for an extended amount of time interactively. This is useful for testing/debugging for an extended amount of time.

```
[cbc@wind ~]$ salloc -c 8 --time=2-00:00:00
salloc: Granted job allocation 33442
[cbc@wind ~]$ srun python analysis.py
[cbc@wind ~]$ exit
salloc: Relinquishing job allocation 33442
[cbc@wind ~]$ salloc -N 2
salloc: Granted job allocation 33443
[cbc@wind ~]$ srun hostname
cn3.nauhpc
cn2.nauhpc
[cbc@wind ~]$ exit
salloc: Relinquishing job allocation 33443
```

Job Parameters

You want	Switches needed
More than one cpu for the job	<code>--cpus-per-task=2</code> , or <code>-c 2</code>
To specify an ordering of your jobs	<code>--dependency=afterok:job_id</code> , or <code>-d job_id</code>
Split up the output, and errors	<code>--output=result.txt --error=error.txt</code>
To run your job at a particular time/day	<code>--begin=16:00 --begin=now+1hour --begin=2010-01-20T12:34:00</code>
Add MPI tasks/ranks to your job	<code>--ntasks=2</code> , or <code>-n 2</code>
To control job failure options	<code>--norequeue --requeue</code>
To receive status email	<code>--mail-type=ALL</code>

Constraints and Resources

You want	Switches needed
To choose a specific node feature (e.g. avx2)	--constraint=avx2
To use a generic resources (e.g. a gpu)	--gres=gpu:tesla:1
To reserve a whole node for yourself	--exclusive
To chose a partition	--partition

Submit the script

```
[cbc@wind ~]$ sbatch jobscript.sh
```

Submitted batch job 85223

- slurm returns a job id for your job that you can use to monitor or modify constraints

Monitoring your job

- `queue`
 - view information about jobs located in the SLURM scheduling queue.
- `queue --start`
- `queue -u login`
- `queue -o "%j %u ... "`
- `queue -p partitionname`
- `queue -S sortfield`
- `queue -t <state> (PD or R)`

Cluster info

- `sinfo`
 - view information about SLURM nodes and partitions.
- `sinfo -N -l`
- `sinfo -R`
 - List reasons for downed nodes and partitions

Monitoring your job

- `sprio`
 - view the factors that comprise a job's scheduling priority
- `sprio -l`
 - list priority of users jobs in pending state
- `sprio -o "%j %u ... "`
- `sprio -w`

Monitoring your job

- `sstat`
 - Display various status information of a running job/step.
- `sstat -j jobid`
- `sstat -o AveCPU,AveRSS`
- Only works with jobs containing steps

Controlling your job

- `scancel`
 - Used to signal jobs or job steps that are under the control of Slurm.
- `scancel -j jobid`
- `scancel -n jobname`
- `scancel -u mylogin`
- `scancel -t pending (only yours)`

Controlling your job

- `scontrol`
 - Used to view and modify Slurm configuration and state.
- `scontrol show job 85224`

Job Accounting

- `sacct`
 - displays accounting data for of your jobs and job steps in the SLURM job accounting log or SLURM database
- `sacct -j jobid -o jobid,elapsed,maxrss`
- `sacct -N nodelist`
- `sacct -u mylogin`
- Try our alias “jobstats”
 - `jobstats`
 - `jobstats -j <jobid>`

Job Accounting

- sshare
 - Tool for listing the shares of associations to a cluster.
- sshare -l : view and compare your fairshare with other accounts
- sshare -a : view all users fairshare
- sshare -A -a <account> : view all members in your account (group)

Account hierarchy

- Your user account belongs to a parent faculty account (group)
- Your user account shares resources that are provided for your group
- Example:
 - coffey
 - cbc
 - mkg52
- View the account structure you belong to with: “sshare -a -A <account>”
- Example:
 - sshare -a -A coffey

Limits on the account (group)

- Limits are in place to prevent intentional or unintentional misuse of resources to ensure quick and fair turn around times on jobs for everyone.
- Groups are limited to a total number of cpu minutes in use at one time: 700,000
- This cpu resource limit mechanism is referred to as: “TRESRunMins”

TRESRunMins Limit

- What the heck is that!?
- A number which limits the total number of remaining cpu minutes which your ***running*** jobs can occupy.
- Enables flexible resource limiting
- Staggers jobs
- Increases cluster utilization
- Leads to more accurate resource requests
- $\text{Sumofjobs}(\text{cpus} * \text{timelimit remaining})$

Examples

- 14400 = 10 jobs, 1 cpu, 1 day in length
- 144000 = 10 jobs, 10 cpu, 1 day in length
- 576000 = 10 jobs, 10 cpu, 5 days in length
- 648000 = 100 jobs, 1 cpu, ½ day in length

Questions?

- Check your groups cpu min usage:
 - sshare -l

Exercise 1

Get to know monsoon and Slurm, on your own.

1. How many nodes make up monsoon?
 - Hint: use “sinfo”
2. How many nodes are in the all partition?
3. How many jobs are currently in the running state ?
 - Hint: use “squeue -t R”
4. How many jobs are currently in the pending state? Why?
 - Hint: use “squeue -t PD”

Exercise 2

- Create a simple job in your home directory
- Example here: `/common/contrib/examples/job_scripts/simplejob.sh` (copy it if you like 😊)
- Name your job: “exercise”
- Name your jobs output: “exercise.out”
- Output should go to `/scratch/<user>/exercise.out`
- Load the module “workshop”
- Run the “date” command
- And additionally, the “secret” command
- Submit your job with sbatch, i.e. “sbatch simplejob.sh”

Exercise 3

- Make your job sleep for 5 minutes (sleep 300)
 - Sleep is a command that creates a process that ... sleeps
- Monitor your job
 - `queue -u your_nuid`
 - `queue -t R`
 - `scontrol show job jobnum`
 - `sacct -j jobnum`
 - Inspect the steps
- Cancel your job
 - `scancel jobnum`

Exercise 4

- Copy job script and edit:
 - `/common/contrib/examples/job_scripts/lazyjob.sh`
- Submit the job, it will take 65 sec to complete
- Use `sstat` and monitor the job
 - `sstat -j <jobid>`
- Review the resources that the job used
 - `jobstats -j <jobid>`
- We are looking for “MaxRSS”, *MaxRSS is the max amount of memory used*
- Edit the job script, reduce the memory being requested in MB and resubmit, edit “`--mem=`”, e.g. `--mem=600`
- Review the resources that the optimized job utilized once again
 - `jobstats -j <jobid>`

- Ok, memory looks good, but notice that the `usercpu` is the same as the elapsed time

$$Usercpu = num\ utilized\ cpus * elapsed\ time$$

- This is because the application we were running only used 1 of the 4 cpus that we requested
- Edit the lazy job script, comment out first `srun` command, and uncomment the second `srun` command.
- Resubmit
- Rerun `jobstats -j <jobid>`, notice now `usercpu` is a multiple times the elapsed time, in this case (4). Because we were allocated 4 cpus, and **used** 4 cpus.

Slurm Arrays

1. Job script is created

```
analysis  
--array=8
```

2. Job script is submitted

```
analysis  
--array=8
```



3. Job is launched with eight instances running in parallel

```
analysis  
123456_1
```

```
analysis  
123456_2
```

```
analysis  
123456_3
```

```
analysis  
123456_4
```

```
analysis  
123456_5
```

```
analysis  
123456_6
```

```
analysis  
123456_7
```

```
analysis  
123456_8
```

Useful environment variables

SLURM_ARRAY_JOB_ID: the job array's ID (parent)

SLURM_ARRAY_TASK_ID: the id of the job array member n (child)

%A

%a

Slurm Arrays Exercise

- From your scratch directory: “/scratch/nauid”
- `tar xvf /common/contrib/examples/bigdata_example.tar`
- `cd bigdata`
- edit the file “job_array.sh” so that it works with your nau id replacing all NAUID with yours
- Submit the script “`sbatch job_array.sh`”
- Run “`queue`”, notice there are 5 jobs running, how did that happen!

MPI Example

- Refer to the MPI example here:
 - `/common/contrib/examples/job_scripts/mpijob.sh`
- Edit it, for your work areas, then experiment:
 - Change number of tasks, nodes ... etc
- Also can run the example like this:
 - `srun --qos=debug -n4 /common/contrib/examples mpi/hellompi`

Keep these tips in mind

- Know the software you are running
- Request resources accurately
- Supply an accurate time limit for your job
- Don't be lazy, it will effect you and your group negatively

Question and Answer

- More info here:
<http://nau.edu/hpc>
- Linux shell help here:
 - <http://linuxcommand.org/tlcl.php>
 - Free book download
- And on the nauhpc listserv
 - nauhpc@lists.nau.edu